JSC270 Assignment 2 Report - *Zihan Guo*

Chosen data: https://archive.ics.uci.edu/ml/machine-learning-databases/adult/adult.data

Background: The *UCI Adult dataset*, also known as the *Census Income dataset*, was extracted from the 1994 U.S. Census database. It contains 48,842 instances and includes various demographic and economic attributes such as *age, workclass, education level, marital status, occupation, race, sex, hours worked per week,* and *income group*. The dataset was originally created for predictive modeling tasks, particularly to determine whether an individual earns more than $50K per year. It is widely used in machine learning, statistical analysis, and labor market research to study income prediction, employment trends, and socioeconomic disparities.

Research Question: How does marital status influence weekly work hours?

Research Details: The goal is to determine whether being married affects the number of hours an individual works per week and whether additional factors such as *education and income* contribute to this relationship. To answer this question, *a linear regression model* was fitted using *hours-per-week* as the dependent variable. The initial model included *marital status* as the sole predictor, where individuals were categorized as either *alone* (if they were divorced, widowed, or never married) or *married* (the rest of non-null status). To refine the model, *education level (education-num)* and *income (income_binary)* were added to control for their potential influence on work hours.

Further improvement was made by incorporating *workclass, occupation, race, relationship status, and native country* to capture broader socioeconomic effects.

Research Process: Using four models to figure out how marital status influence the working

time. Each model builds upon the previous one, improving explanatory power and reducing prediction error, while also revealing potential model limitations. In Model 1(Appendix 1&5), *marital status* is the only predictor, showing that married individuals work *5.22 more hours* per week than those who are not married. This relationship is statistically significant (p-value = 0.000), but the low R-squared (0.044) suggests that marital status alone explains only a small portion of the variance in work hours. The high RMSE (12.07) indicates significant prediction errors, implying missing influential factors. Model 2(Appendix 2&5) adds *education (education-num)*, reducing the marital status effect to 4.96 hours, meaning education partly explains work hour differences. Each additional year of education *increases work hours by 0.64*, and the model fit improves (R-squared = 0.062, RMSE = 11.96), though errors remain high. Model 3(Appendix 3&5) incorporates income (income_binary), further reducing the marital status effect, confirming that income explains part of the relationship. The model's predictive power improves (R-squared = 0.075, RMSE = 11.87), but residual errors remain large. Model 4[improving model](Appendix 4&5) expands with additional predictors (*workclass, occupation, race, relationship status, and native country*), yielding the best fit (*R-squared increases, RMSE drops to 11.07*). However, multicollinearity issues suggest that some predictors are highly correlated, potentially distorting coefficient estimates.

Despite improvements, RMSE remains high, likely due to high variability in work hours, missing key predictors (job type, industry, household responsibilities), multicollinearity, and potential non-linearity. To reduce errors, we should remove multicollinear variables (using VIF), apply log transformations to skewed variables, introduce interaction terms, and explore non-linear models like decision trees or random forests to improve predictive accuracy.

Appendix:

Appendix 1:

Model 1: Marital Status Only

OLS Regression Results

Dep. Variable:	Q('hours-per-week')	R-squared:	0.044
Model:	OLS	Adj. R-squared:	0.044
Method:	Least Squares	F-statistic:	1516.
Date:	Tue, 18 Feb 2025	Prob (F-statistic):	4.94e-324
Time:	05:00:06	Log-Likelihood:	-1.2730e+05
No. Observations:	32561	AIC:	2.546e+05
Df Residuals:	32559	BIC:	2.546e+05
Df Model:	1		
Covariance Type:	nonrobust		

	coef	std err	t	P>|t|	[0.025	0.975]
Intercept	37.9679	0.092	411.880	0.000	37.787	38.149
Q('current_married_married')[T.True]	5.2157	0.134	38.933	0.000	4.953	5.478

Omnibus:	2686.362	Durbin-Watson:	2.016
Prob(Omnibus):	0.000	Jarque-Bera (JB):	13049.893
Skew:	0.257	Prob(JB):	0.00
Kurtosis:	6.059	Cond. No.	2.56

Notes:
[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

Appendix 2:

Model 2: Marital Status and Education

OLS Regression Results

Dep. Variable:	Q('hours-per-week')	R-squared:	0.062
Model:	OLS	Adj. R-squared:	0.062
Method:	Least Squares	F-statistic:	1074.
Date:	Tue, 18 Feb 2025	Prob (F-statistic):	0.00
Time:	05:00:06	Log-Likelihood:	-1.2700e+05
No. Observations:	32561	AIC:	2.540e+05
Df Residuals:	32558	BIC:	2.540e+05
Df Model:	2		
Covariance Type:	nonrobust		

	coef	std err	t	P>|t|	[0.025	0.975]
Intercept	31.6823	0.271	116.739	0.000	31.150	32.214
Q('current_married_married')[T.True]	4.9594	0.133	37.247	0.000	4.698	5.220
Q('education-num')	0.6356	0.026	24.595	0.000	0.585	0.686

Omnibus:	2804.860	Durbin-Watson:	2.016
Prob(Omnibus):	0.000	Jarque-Bera (JB):	14359.168
Skew:	0.260	Prob(JB):	0.00
Kurtosis:	6.211	Cond. No.	43.2

Notes:
[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

Appendix 3:

Model 3: Marital Status, Education, and Income

OLS Regression Results

Dep. Variable:	Q('hours-per-week')	R-squared:	0.075
Model:	OLS	Adj. R-squared:	0.075
Method:	Least Squares	F-statistic:	883.5
Date:	Tue, 18 Feb 2025	Prob (F-statistic):	0.00
Time:	05:00:06	Log-Likelihood:	-1.2677e+05
No. Observations:	32561	AIC:	2.535e+05
Df Residuals:	32557	BIC:	2.536e+05
Df Model:	3		
Covariance Type:	nonrobust		

	coef	std err	t	P>|t|	[0.025	0.975]
Intercept	33.3886	0.281	118.947	0.000	32.838	33.939
Q('current_married_married')[T.True]	3.5734	0.147	24.337	0.000	3.286	3.861
Q('education-num')	0.4374	0.027	16.060	0.000	0.384	0.491
income_binary	3.9359	0.181	21.694	0.000	3.580	4.291

Omnibus:	2929.612	Durbin-Watson:	2.014
Prob(Omnibus):	0.000	Jarque-Bera (JB):	14973.351
Skew:	0.290	Prob(JB):	0.00
Kurtosis:	6.271	Cond. No.	46.7

Notes:
[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

Appendix 4:

Model 4: Extended Model with Additional Predictors
OLS Regression Results

Dep. Variable:	Q('hours-per-week')	R-squared:	0.195
Model:	OLS	Adj. R-squared:	0.193
Method:	Least Squares	F-statistic:	85.57
Date:	Tue, 18 Feb 2025	Prob (F-statistic):	0.00
Time:	05:00:08	Log-Likelihood:	-1.2451e+05
No. Observations:	32561	AIC:	2.492e+05
Df Residuals:	32468	BIC:	2.500e+05
Df Model:	92		
Covariance Type:	nonrobust		

	coef	std err	t	P>|t|	[0.025	0.975]
Intercept	35.0575	1.064	32.955	0.000	32.972	37.143
workclass[T.Federal-gov]	4.0563	0.443	9.166	0.000	3.189	4.924
workclass[T.Local-gov]	4.2414	0.380	11.157	0.000	3.496	4.987
workclass[T.Never-worked]	-1.4184	4.204	-0.337	0.736	-9.659	6.822
workclass[T.Private]	4.1511	0.304	13.669	0.000	3.556	4.746
workclass[T.Self-emp-inc]	9.1262	0.441	20.698	0.000	8.262	9.990
workclass[T.Self-emp-not-inc]	5.8580	0.366	15.989	0.000	5.140	6.576
workclass[T.State-gov]	2.1655	0.420	5.161	0.000	1.343	2.988
workclass[T.Without-pay]	-2.3128	2.841	-0.814	0.416	-7.880	3.255
education[T.11th]	-2.2977	0.488	-4.711	0.000	-3.254	-1.342
education[T.12th]	-0.8704	0.647	-1.346	0.178	-2.138	0.397
education[T.1st-4th]	0.2397	0.962	0.249	0.803	-1.645	2.125
education[T.5th-6th]	0.5989	0.747	0.802	0.423	-0.865	2.063
education[T.7th-8th]	1.2168	0.573	2.123	0.034	0.094	2.340
education[T.9th]	0.6341	0.612	1.036	0.300	-0.566	1.834
education[T.Assoc-acdm]	1.8574	0.505	3.679	0.000	0.868	2.847
education[T.Assoc-voc]	2.5792	0.477	5.412	0.000	1.645	3.513
education[T.Bachelors]	2.5190	0.412	6.108	0.000	1.711	3.327
education[T.Doctorate]	6.2281	0.695	8.962	0.000	4.866	7.590
education[T.HS-grad]	2.0927	0.381	5.490	0.000	1.346	2.840

...

native[T.Portugal]	2.1859	1.893	1.155	0.248	-1.525	5.897
native[T.Puerto-Rico]	-0.6207	1.146	-0.541	0.588	-2.867	1.626
native[T.Scotland]	1.4419	3.238	0.445	0.656	-4.906	7.789
native[T.South]	2.2565	1.396	1.617	0.106	-0.479	4.992
native[T.Taiwan]	-3.5237	1.679	-2.098	0.036	-6.815	-0.232
native[T.Thailand]	4.8495	2.687	1.805	0.071	-0.417	10.116
native[T.Trinadad&Tobago]	-1.4177	2.593	-0.547	0.585	-6.499	3.664
native[T.United-States]	-0.4134	0.473	-0.875	0.382	-1.340	0.513
native[T.Vietnam]	-1.9585	1.500	-1.306	0.192	-4.898	0.981
native[T.Yugoslavia]	1.8702	2.814	0.665	0.506	-3.646	7.386
current_married_married[T.True]	1.4838	0.453	3.277	0.001	0.596	2.371
age	-0.0981	0.005	-18.190	0.000	-0.109	-0.088
fnlwgt	-1.978e-06	6e-07	-3.298	0.001	-3.15e-06	-8.03e-07
capital	2.805e-05	8.66e-06	3.238	0.001	1.11e-05	4.5e-05
sex_binary	1.3504	0.092	14.724	0.000	1.171	1.530
income_binary	3.1175	0.179	17.408	0.000	2.767	3.469

Omnibus:	3673.644	Durbin-Watson:	2.015
Prob(Omnibus):	0.000	Jarque-Bera (JB):	18168.532
Skew:	0.447	Prob(JB):	0.00
Kurtosis:	6.549	Cond. No.	8.74e+18

Appendix 5:

```
Model 1: Marital Status Only:
Residual Standard Error (RSE): 12.0699
Root Mean Squared Error (RMSE): 12.0695

Model 2: Marital Status and Education:
Residual Standard Error (RSE): 11.9595
Root Mean Squared Error (RMSE): 11.9589

Model 3: Marital Status, Education, and Income:
Residual Standard Error (RSE): 11.8741
Root Mean Squared Error (RMSE): 11.8734

Model 4: Extended Model with Additional Predictors:
Residual Standard Error (RSE): 11.0930
Root Mean Squared Error (RMSE): 11.0771
```